

# A DESCRIPTIVE TITLE, NOT TOO GENERAL, NOT TOO LONG

*Markus Püschel*

Department of Computer Science  
ETH Zürich  
Zürich, Switzerland

The hard page limit is 6 pages in this style. Do not reduce font size or use other tricks to squeeze. This pdf is formatted in the American letter format, so the spacing may look a bit strange when printed out.

## ABSTRACT

Describe in concise words what you do, why you do it (not necessarily in this order), and the main result. The abstract has to be self-contained and readable for a person in the general area. You should write the abstract last.

## 1. INTRODUCTION

**Motivation.** Breadth-first search is a graph traversing algorithm which is used in many applications. Some examples are finding connected components, determining a shortest-path in a navigation system or computing the maximum flow in a network. As real-world graphs can get very large, for example the US road network graph contains roughly 24 million vertices and 57 million edges, the need for fast implementations is high.

In this paper, we will discuss multiple approaches and compare their performance on three different platforms.

**Related work.** In this section we will give a brief overview of related work. In the first part we will discuss the work covering the top-down-bottom-up hybrid approach and in the second part an approach that focuses on avoiding atomic operations.

[1] [2]

Berrendorf [3] describes a technique to avoid atomic operations in a generalized scenario. The scenario is given as an if-statement followed by some operations that change a state, where multiple threads might execute the predicate and execute the operations afterwards. The operations need to change the state to the same value if executed multiple times otherwise there exists a race condition, i.e. the change of the distance of a visited vertex to the value of the level or the addition of a vertex to the next frontier. The trade-off is

---

The author thanks Jelena Kovacevic. This paper is a modified version of the template she used in her class.

that doing a BFS this way can result in additional work any unvisited vertex may get added multiple times.

## 2. BACKGROUND: WHATEVER THE BACKGROUND IS

Give a short, self-contained summary of necessary background information. For example, assume you present an implementation of sorting algorithms. You could organize into sorting definition, algorithms considered, and asymptotic runtime statements. The goal of the background section is to make the paper self-contained for an audience as large as possible. As in every section you start with a very brief overview of the section. Here it could be as follows: In this section we formally define the sorting problem we consider and introduce the algorithms we use including a cost analysis.

**Sorting.** Precisely define sorting problem you consider.

**Sorting algorithms.** Explain the algorithm you use including their costs.

As an aside, don't talk about "the complexity of the algorithm." It's incorrect, problems have a complexity, not algorithms.

## 3. YOUR PROPOSED METHOD

Now comes the "beef" of the report, where you explain what you did. Again, organize it in paragraphs with titles. As in every section you start with a very brief overview of the section.

In this section, structure is very important so one can follow the technical content.


Mention and cite any external resources that you used including libraries or other code.

## 4. EXPERIMENTAL RESULTS

Here you evaluate your work using experiments. You start again with a very short summary of the section. The typical structure follows.

DFT (single precision) on Intel Core i7 (4 cores)

Performance [Gflop/s] vs. input size



**Fig. 1.** Performance of four single precision implementations of the discrete Fourier transform. The operations count is roughly the same. The labels in this plot are maybe a little bit too small.

**Experimental setup.** We run experiments on three different platforms.

The first platform is the EULER cluster which is operated by the HPC Group of ETH. We had access to one node with a 12-core Intel Xeon E5-2697v2 processors (2.7 GHz nominal, 3.0-3.5 GHz peak). It supports hyper-threading, so we ran our algorithms with up to 24 threads. On this platform, we used the gcc compiler with the -O2 flag.

The next platform we ran our algorithms on is the Xeon Phi provided through the class.

Lastly, we also used an AMD FX-8350 (4GHz x8, 8Gb, W2k3).

**Results.** Next divide the experiments into classes, one paragraph for each. In each class of experiments you typically pursue one questions that then is answered by a suitable plot or plots. For example, first you may want to investigate the performance behavior with changing input size, then how your code compares to external benchmarks.

For some tips on benchmarking including how to create a decent viewgraph see pages 22–27 in

#### Comments:

- Create very readable, attractive plots (do 1 column, not 2 column plots for this report) with readable font size. However, the font size should also not be too large; typically it is smaller than the text font size. An example is in Fig. 1 (of course you can have a different style).
- Every plot answers a question. You state this question and extract the answer from the plot in its discussion.
- Every plot should be referenced and discussed.

## 5. CONCLUSIONS

Here you need to summarize what you did and why this is important. *Do not take the abstract* and put it in the past tense. Remember, now the reader has (hopefully) read the report, so it is a very different situation from the abstract. Try to highlight important results and say the things you really want to get across such as high-level statements (e.g., we believe that .... is the right approach to .... Even though we only considered x, the .... technique should be applicable ....) You can also formulate next steps if you want. Be brief. After the conclusions there are only the references.

## 6. FURTHER COMMENTS

Here we provide some further tips.

#### Further general guidelines.

- For short papers, to save space, I use paragraph titles instead of subsections, as shown in the introduction.
- It is generally a good idea to break sections into such smaller units for readability and since it helps you to (visually) structure the story.
- The above section titles should be adapted to more precisely reflect what you do.
- Each section should be started with a very short summary of what the reader can expect in this section. Nothing more awkward as when the story starts and one does not know what the direction is or the goal.
- Make sure you define every acronym you use, no matter how convinced you are the reader knows it.
- Always spell-check before you submit (to us in this case).
- Be picky. When writing a paper you should always strive for very high quality. Many people may read it and the quality makes a big difference. In this class, the quality is part of the grade.
- Books helping you to write better:
- Conversion to pdf (latex users only):  

```
dvips -o conference.ps -t letter -Ppdf -G0 conference.dvi
```

and then  

```
ps2pdf conference.ps
```

**Graphics.** For plots that are not images *never* generate the bitmap formats jpeg, gif, bmp, tif. Use eps, which means encapsulate postscript. It is scalable since it is a vector graphic description of your graph. E.g., from Matlab, you can export to eps.

The format pdf is also fine for plots (you need pdf<sub>l</sub>atex then), but only if the plot was never before in the format jpeg, gif, bmp, tif.

## 7. REFERENCES

- [1] Y. Yasui, K. Fujisawa, and K. Goto, “Numa-optimized parallel breadth-first search on multicore single-node system,” in *Big Data, 2013 IEEE International Conference on*, Oct 2013, pp. 394–402.
- [2] Scott Beamer, Krste Asanovi, and David A Patterson, “Searching for a parent instead of fighting over children: A fast breadth-first search implementation for graph500,” 2011.
- [3] R. Berrendorf, “A technique to avoid atomic operations on large shared memory parallel systems,” *International Journal on Advances in Software*, vol. 7, no. 1 & 2, pp. 197–210, 2014.